

Machine_Learning_Medical_Insurance

December 10, 2022

1 A Machine Learning Example for the Insurance Sector

Dataset: <https://www.kaggle.com/datasets/mirichoi0218/insurance>

```
[1]: import pandas as pd
import numpy as np
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

sns.set_style("darkgrid")
matplotlib.rcParams['font.size'] = 14
matplotlib.rcParams['figure.figsize'] = (15, 5)
matplotlib.rcParams['figure.facecolor'] = '#00000000'

import warnings
warnings.simplefilter(action='ignore')
```

2 Importing Data

```
[2]: insurance_df = pd.read_csv('./insurance.csv')
insurance_df.head()
```

```
[2]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
[3]: insurance_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -

```

```

0   age      1338 non-null   int64
1   sex      1338 non-null   object
2   bmi      1338 non-null   float64
3   children 1338 non-null   int64
4   smoker   1338 non-null   object
5   region   1338 non-null   object
6   charges  1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB

```

```
[4]: insurance_df[insurance_df.duplicated(keep=False)]
```

```

[4]:      age  sex  bmi  children  smoker  region  charges
195   19  male  30.59          0     no  northwest  1639.5631
581   19  male  30.59          0     no  northwest  1639.5631

```

There is one duplicated value. It is better to drop it.

```
[5]: insurance_df
```

```

[5]:      age  sex  bmi  children  smoker  region  charges
0      19  female  27.900          0    yes  southwest  16884.92400
1      18   male  33.770          1     no  southeast  1725.55230
2      28   male  33.000          3     no  southeast  4449.46200
3      33   male  22.705          0     no  northwest  21984.47061
4      32   male  28.880          0     no  northwest  3866.85520
...  ...  ...  ...  ...  ...  ...
1333   50   male  30.970          3     no  northwest  10600.54830
1334   18  female  31.920          0     no  northeast  2205.98080
1335   18  female  36.850          0     no  southeast  1629.83350
1336   21  female  25.800          0     no  southwest  2007.94500
1337   61  female  29.070          0    yes  northwest  29141.36030

```

[1338 rows x 7 columns]

```

[6]: insurance_df_no_dupl = insurance_df.drop_duplicates()
insurance_df_no_dupl

```

```

[6]:      age  sex  bmi  children  smoker  region  charges
0      19  female  27.900          0    yes  southwest  16884.92400
1      18   male  33.770          1     no  southeast  1725.55230
2      28   male  33.000          3     no  southeast  4449.46200
3      33   male  22.705          0     no  northwest  21984.47061
4      32   male  28.880          0     no  northwest  3866.85520
...  ...  ...  ...  ...  ...  ...
1333   50   male  30.970          3     no  northwest  10600.54830
1334   18  female  31.920          0     no  northeast  2205.98080
1335   18  female  36.850          0     no  southeast  1629.83350

```

```
1336    21  female  25.800          0    no  southwest    2007.94500
1337    61  female  29.070          0   yes  northwest   29141.36030
```


[1337 rows x 7 columns]

3 Exploratory Data Analysis

3.1 Exploring the Target Variable 'charges'

Before starting with the other variables it is important to analyze the distribution of the 'charges' feature.

```
[7]: sns.histplot(data=insurance_df_no_dupl, x='charges', bins=50 ,kde=True);
```



```
[8]: insurance_df_no_dupl.charges.describe()
```

```
[8]: count      1337.000000
     mean      13279.121487
     std       12110.359656
     min       1121.873900
     25%       4746.344000
     50%       9386.161300
     75%      16657.717450
     max       63770.428010
     Name: charges, dtype: float64
```

The distribution is right-skewed.


The majority of the fees are concentrated in the range between \$0 and \$10000.

3.2 Exploring the Variable Age

```
[9]: insurance_df_no_dupl.age.describe()
```

```
[9]: count    1337.000000
     mean      39.222139
     std       14.044333
     min       18.000000
     25%       27.000000
     50%       39.000000
     75%       51.000000
     max       64.000000
     Name: age, dtype: float64
```

```
[10]: sns.boxplot(data=insurance_df_no_dupl, x='age');
```



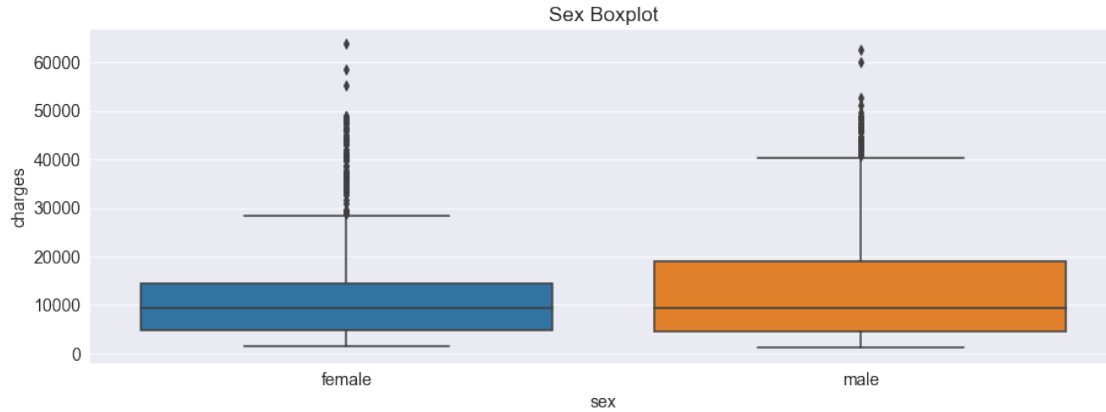
3.3 Exploring the Sex Variable

```
[11]: insurance_df_no_dupl.sex.value_counts()
```

```
[11]: male        675
     female      662
     Name: sex, dtype: int64
```

The distribution of sex is balanced. What is the amount of charges paid by sex?

```
[12]: sns.boxplot(data=insurance_df_no_dupl, x='sex', y='charges')
     plt.title('Sex Boxplot');
```



Males seem to pay quite more in charges than females.

```
[13]: insurance_df_no_dupl.groupby('sex')['charges'].mean()
```

```
[13]: sex
      female    12569.578844
      male     13974.998864
      Name: charges, dtype: float64
```

```
[14]: females_charges = insurance_df_no_dupl.groupby('sex')['charges'].mean()[0]
      male_charges = insurance_df_no_dupl.groupby('sex')['charges'].mean()[1]
      male_charges-females_charges
```

```
[14]: 1405.4200199276147
```

Indeed, on average, males pay \$1405.42 more.

A more detailed description of the variable “charges” described by gender.

```
[15]: insurance_df_no_dupl[insurance_df_no_dupl.sex=='male'].charges.describe()
```

```
[15]: count      675.000000
      mean      13974.998864
      std       12971.958663
      min       1121.873900
      25%       4654.022675
      50%       9377.904700
      75%      19006.685500
      max      62592.873090
      Name: charges, dtype: float64
```

```
[16]: insurance_df_no_dupl[insurance_df_no_dupl.sex=='female'].charges.describe()
```

```
[16]: count      662.000000
      mean      12569.578844
      std       11128.703801
      min       1607.510100
      25%       4885.158700
      50%       9412.962500
      75%      14454.691825
      max       63770.428010
      Name: charges, dtype: float64
```


3.4 Exploring the BMI Variable

```
[17]: insurance_df_no_dupl.bmi.describe()
```

```
[17]: count      1337.000000
      mean       30.663452
      std        6.100468
      min       15.960000
      25%       26.290000
      50%       30.400000
      75%       34.700000
      max       53.130000
      Name: bmi, dtype: float64
```

Is there a correlation with “charges” feature?


```
[18]: sns.scatterplot(data=insurance_df_no_dupl, x='bmi', y='charges', hue='sex')
      plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0);
```



It does not seem to be a strong relationship, but it is positive for sure; moreover there are not two well defined clusters based on sex.

Checking the distribution of the variable BMI.

```
[19]: sns.histplot(data=insurance_df_no_dupl, x='bmi', bins=50, kde=True);
```




The “bmi” variable seems to be normally distributed.

3.5 Exploring the Children Variable

```
[20]: insurance_df_no_dupl.children.describe()
```

```
[20]: count      1337.000000
      mean         1.095737
      std          1.205571
      min          0.000000
      25%          0.000000
      50%          1.000000
      75%          2.000000
      max          5.000000
      Name: children, dtype: float64
```

```
[21]: sns.boxplot(data=insurance_df_no_dupl, x='children', y='charges');
```



```
[22]: insurance_df_no_dupl.groupby('children')['charges'].mean()
```

```
[22]: children
0    12384.695344
1    12731.171832
2    15073.563734
3    15355.318367
4    13850.656311
5     8786.035247
Name: charges, dtype: float64
```

On average, who has 2 or 3 children pays spends more money in charges.


3.6 Exploring the Smoking Variable

```
[23]: insurance_df_no_dupl.smoker.value_counts()
```

```
[23]: no      1063
yes      274
Name: smoker, dtype: int64
```

There are not many smokers in the dataset.

```
[24]: sns.boxplot(data=insurance_df_no_dupl, x='smoker', y='charges');
```



In general, smokers pay much more in insurance charges than no-smokers.

```
[25]: insurance_df_no_dupl[insurance_df_no_dupl.smoker=='yes'].charges.describe()
```

```
[25]: count      274.000000
mean      32050.231832
std       11541.547176
min       12829.455100
25%       20826.244213
```



```
50%      34456.348450
75%      41019.207275
max       63770.428010
Name: charges, dtype: float64
```

```
[26]: insurance_df_no_dupl[insurance_df_no_dupl.smoker=='no'].charges.describe()
```

```
[26]: count      1063.000000
      mean      8440.660307
      std      5992.973800
      min      1121.873900
      25%      3988.883500
      50%      7345.726600
      75%     11363.019100
      max      36910.608030
      Name: charges, dtype: float64
```

What is, on average, the amount paid by smokers and non-smokers?

```
[27]: insurance_df_no_dupl.groupby('smoker')['charges'].mean()
```

```
[27]: smoker
      no      8440.660307
      yes    32050.231832
      Name: charges, dtype: float64
```

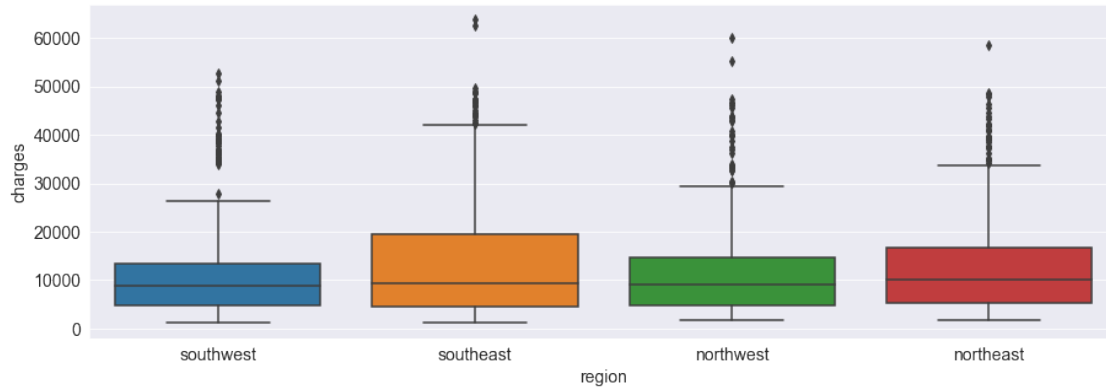
3.7 Exploring the Region Variable

```
[28]: insurance_df_no_dupl.region.value_counts()
```

```
[28]: southeast      364
      southwest     325
      northwest     324
      northeast     324
      Name: region, dtype: int64
```

Regions are almost perfectly balanced, only the southeast region appears more times.

```
[29]: sns.boxplot(data=insurance_df_no_dupl, x='region', y='charges');
```



It does not seem to be a large difference between them. Only the southeast region has more outliers.

4 Feature Engineering

In this section are performed all the necessary steps to transform variables that will be used in the machine learning model.

4.1 Defining Inputs and Target

```
[30]: input_df = insurance_df_no_dupl.drop(columns='charges')
      target_df = insurance_df_no_dupl.charges
```

Showing the results

```
[31]: input_df
```

```
[31]:
```

	age	sex	bmi	children	smoker	region
0	19	female	27.900	0	yes	southwest
1	18	male	33.770	1	no	southeast
2	28	male	33.000	3	no	southeast
3	33	male	22.705	0	no	northwest
4	32	male	28.880	0	no	northwest
...
1333	50	male	30.970	3	no	northwest
1334	18	female	31.920	0	no	northeast
1335	18	female	36.850	0	no	southeast
1336	21	female	25.800	0	no	southwest
1337	61	female	29.070	0	yes	northwest

[1337 rows x 6 columns]

```
[32]: target_df
```

```
[32]: 0      16884.92400
      1      1725.55230
      2      4449.46200
      3     21984.47061
      4      3866.85520
      ...
     1333    10600.54830
     1334     2205.98080
     1335     1629.83350
     1336     2007.94500
     1337    29141.36030
      Name: charges, Length: 1337, dtype: float64
```

4.2 Creating a Variable to Cluster the BMI

It will be added a new feature to cluster observations on their BMI values.

These are the different ranges:

- If BMI is less than 18.5, it falls within the underweight range.
- If BMI is 18.5 to 24.9, it falls within the normal or Healthy Weight range.
- If BMI is 25.0 to 29.9, it falls within the overweight range.
- If BMI is 30.0 or higher, it falls within the obese range.

```
[33]: def bmi_estimator(column):
      if column < 18.5:
          return 'underweight'
      elif (column >= 18.5) and (column<=24.9):
          return 'healthy weight'
      elif (column>=25) and (column<=29.9):
          return 'overweight'
      else:
          return 'obese'
```

```
[34]: input_df['bmi_class']=input_df.bmi.map(bmi_estimator)
      input_df
```

```
[34]:   age  sex  bmi  children  smoker  region  bmi_class
0    19 female 27.900         0    yes southwest  overweight
1    18  male 33.770         1    no  southeast    obese
2    28  male 33.000         3    no  southeast    obese
3    33  male 22.705         0    no northwest healthy weight
4    32  male 28.880         0    no northwest  overweight
...  ...  ...  ...  ...  ...  ...  ...
1333 50  male 30.970         3    no northwest    obese
1334 18 female 31.920         0    no northeast    obese
1335 18 female 36.850         0    no southeast    obese
1336 21 female 25.800         0    no southwest  overweight
1337 61 female 29.070         0    yes northwest  overweight
```

[1337 rows x 7 columns]

4.3 Encoding Sex and Smoking Variables

```
[35]: sex_dict = {'male':0, 'female':1}
      smoker_dict = {'no':0, 'yes':1}

      input_df['sex'] = input_df.sex.map(sex_dict)
      input_df['smoker'] = input_df.smoker.map(smoker_dict)
```

4.4 Encoding the Region Variable

It is used the one-hot encoding procedure.

```
[36]: columns_to_encode = ['region', 'bmi_class']
```

```
[37]: from sklearn.preprocessing import OneHotEncoder

      encoder = OneHotEncoder(sparse=False, handle_unknown='ignore').
      ↪fit(input_df[columns_to_encode])

      encoded_cols = list(encoder.get_feature_names_out(columns_to_encode))

      input_df[encoded_cols] = encoder.transform(input_df[columns_to_encode])
```

```
[38]: input_df = input_df.drop(columns='region')
      input_df
```

```
[38]:
```

	age	sex	bmi	children	smoker	bmi_class	region_northeast \
0	19	1	27.900	0	1	overweight	0.0
1	18	0	33.770	1	0	obese	0.0
2	28	0	33.000	3	0	obese	0.0
3	33	0	22.705	0	0	healthy weight	0.0
4	32	0	28.880	0	0	overweight	0.0
...
1333	50	0	30.970	3	0	obese	0.0
1334	18	1	31.920	0	0	obese	1.0
1335	18	1	36.850	0	0	obese	0.0
1336	21	1	25.800	0	0	overweight	0.0
1337	61	1	29.070	0	1	overweight	0.0

	region_northwest	region_southeast	region_southwest \
0	0.0	0.0	1.0
1	0.0	1.0	0.0
2	0.0	1.0	0.0
3	1.0	0.0	0.0
4	1.0	0.0	0.0

...
1333	1.0	0.0	0.0
1334	0.0	0.0	0.0
1335	0.0	1.0	0.0
1336	0.0	0.0	1.0
1337	1.0	0.0	0.0

	bmi_class_healthy	weight	bmi_class_obese	bmi_class_overweight	\
0		0.0	0.0	1.0	
1		0.0	1.0	0.0	
2		0.0	1.0	0.0	
3		1.0	0.0	0.0	
4		0.0	0.0	1.0	
...		
1333		0.0	1.0	0.0	
1334		0.0	1.0	0.0	
1335		0.0	1.0	0.0	
1336		0.0	0.0	1.0	
1337		0.0	0.0	1.0	

	bmi_class_underweight
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
1333	0.0
1334	0.0
1335	0.0
1336	0.0
1337	0.0

[1337 rows x 14 columns]

4.5 Scaling Values

Values will be scaled by using MinMax Scaler.

```
[39]: columns_to_scale = ['age', 'bmi', 'children']
input_df[columns_to_scale].head()
```

```
[39]:   age    bmi  children
0   19  27.900         0
1   18  33.770         1
2   28  33.000         3
3   33  22.705         0
```

4 32 28.880 0

```
[40]: from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler().fit(input_df[columns_to_scale])
```

```
input_df[columns_to_scale] = scaler.transform(input_df[columns_to_scale])
```

```
[41]: input_df = input_df.drop(columns='bmi_class')
input_df.head()
```

```
[41]:
```

	age	sex	bmi	children	smoker	region_northeast	\
0	0.021739	1	0.321227	0.0	1	0.0	
1	0.000000	0	0.479150	0.2	0	0.0	
2	0.217391	0	0.458434	0.6	0	0.0	
3	0.326087	0	0.181464	0.0	0	0.0	
4	0.304348	0	0.347592	0.0	0	0.0	

	region_northwest	region_southeast	region_southwest	\
0	0.0	0.0	1.0	
1	0.0	1.0	0.0	
2	0.0	1.0	0.0	
3	1.0	0.0	0.0	
4	1.0	0.0	0.0	

	bmi_class_healthy	weight	bmi_class_obese	bmi_class_overweight	\
0	0.0	0.0	0.0	1.0	
1	0.0	0.0	1.0	0.0	
2	0.0	0.0	1.0	0.0	
3	1.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	1.0	

	bmi_class_underweight
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

```
[42]: input_df.describe().loc[['min', 'max']]
```

```
[42]:
```

	age	sex	bmi	children	smoker	region_northeast	region_northwest	\
min	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
max	1.0	1.0	1.0	1.0	1.0	1.0	1.0	

	region_southeast	region_southwest	bmi_class_healthy	weight	\
min	0.0	0.0	0.0	0.0	

max	1.0	1.0	1.0
	bmi_class_obese	bmi_class_overweight	bmi_class_underweight
min	0.0	0.0	0.0
max	1.0	1.0	1.0

All the variables range now between 0 and 1.

4.6 Splitting the Input Data Frame into a Train and Validation One

```
[43]: from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(input_df, target_df,
↪test_size=0.2, random_state=42)
```

Showing the result.

```
[44]: X_train
```

```
[44]:
```

	age	sex	bmi	children	smoker	region_northeast	\
1114	0.108696	0	0.230024	0.0	0	1.0	
968	0.065217	0	0.263250	0.4	0	1.0	
599	0.739130	1	0.580172	0.4	0	0.0	
170	0.978261	0	0.686306	0.0	0	0.0	
275	0.630435	1	0.286252	0.4	0	1.0	
...	
1096	0.717391	1	0.511165	0.4	1	1.0	
1131	0.195652	0	0.805488	0.4	0	0.0	
1295	0.043478	0	0.162497	0.2	0	0.0	
861	0.434783	1	0.323917	0.6	0	0.0	
1127	0.369565	1	0.535378	0.4	0	0.0	

	region_northwest	region_southeast	region_southwest	\
1114	0.0	0.0	0.0	
968	0.0	0.0	0.0	
599	1.0	0.0	0.0	
170	0.0	1.0	0.0	
275	0.0	0.0	0.0	
...	
1096	0.0	0.0	0.0	
1131	0.0	0.0	1.0	
1295	0.0	0.0	1.0	
861	0.0	0.0	1.0	
1127	0.0	1.0	0.0	

	bmi_class_healthy weight	bmi_class_obese	bmi_class_overweight	\
1114	1.0	0.0	0.0	
968	0.0	0.0	1.0	

599	0.0	1.0	0.0
170	0.0	1.0	0.0
275	0.0	0.0	1.0
...
1096	0.0	1.0	0.0
1131	0.0	1.0	0.0
1295	1.0	0.0	0.0
861	0.0	0.0	1.0
1127	0.0	1.0	0.0

	bmi_class_underweight
1114	0.0
968	0.0
599	0.0
170	0.0
275	0.0
...	...
1096	0.0
1131	0.0
1295	0.0
861	0.0
1127	0.0

[1069 rows x 13 columns]

[45]: X_val

	age	sex	bmi	children	smoker	region_northeast	\
900	0.673913	0	0.176352	0.0	0	1.0	
1064	0.239130	1	0.259349	0.8	0	0.0	
1256	0.717391	1	0.549502	0.6	0	0.0	
298	0.282609	0	0.495830	0.6	1	0.0	
237	0.282609	0	0.603444	0.4	0	0.0	
...	
534	1.000000	0	0.659672	0.0	0	0.0	
542	0.978261	1	0.547215	0.0	0	0.0	
760	0.086957	1	0.500942	0.4	0	1.0	
1284	0.934783	0	0.547215	0.2	1	0.0	
1285	0.630435	1	0.224913	0.0	0	1.0	

	region_northwest	region_southeast	region_southwest	\
900	0.0	0.0	0.0	
1064	0.0	0.0	1.0	
1256	1.0	0.0	0.0	
298	1.0	0.0	0.0	
237	0.0	1.0	0.0	
...	

534	0.0	1.0	0.0
542	0.0	1.0	0.0
760	0.0	0.0	0.0
1284	0.0	0.0	1.0
1285	0.0	0.0	0.0

	bmi_class_healthy	weight	bmi_class_obese	bmi_class_overweight	\
900		1.0	0.0	0.0	
1064		0.0	0.0	1.0	
1256		0.0	1.0	0.0	
298		0.0	1.0	0.0	
237		0.0	1.0	0.0	
...		
534		0.0	1.0	0.0	
542		0.0	1.0	0.0	
760		0.0	1.0	0.0	
1284		0.0	1.0	0.0	
1285		1.0	0.0	0.0	

	bmi_class_underweight
900	0.0
1064	0.0
1256	0.0
298	0.0
237	0.0
...	...
534	0.0
542	0.0
760	0.0
1284	0.0
1285	0.0

[268 rows x 13 columns]

```
[46]: y_train
```

```
[46]: 1114    2396.09590
      968    3279.86855
      599   33471.97189
      170   13405.39030
      275    9715.84100
      ...
      1096   44641.19740
      1131    3693.42800
      1295    1964.78000
      861    7151.09200
      1127    5836.52040
```

Name: charges, Length: 1069, dtype: float64

```
[47]: y_val
```

```
[47]: 900      8688.85885
      1064      5708.86700
      1256     11436.73815
      298     38746.35510
      237      4463.20510
      ...
      534     13831.11520
      542     13887.20400
      760      3925.75820
      1284    47403.88000
      1285     8534.67180
      Name: charges, Length: 268, dtype: float64
```

5 Creating the Model

It will be used the XGBRegressor model

```
[48]: from xgboost import XGBRegressor
```

```
[49]: model = XGBRegressor(n_jobs=-1, n_estimators=1000, early_stopping_rounds=50,
    ↪random_state=42)
      model.fit(X_train, y_train, eval_set=[(X_train, y_train), (X_val, y_val)])
```

```
[0]      validation_0-rmse:12740.82085      validation_1-rmse:14354.43509
[1]      validation_0-rmse:9486.62138      validation_1-rmse:10616.10467
[2]      validation_0-rmse:7287.93828      validation_1-rmse:8240.67071
[3]      validation_0-rmse:5864.49266      validation_1-rmse:6655.85239
[4]      validation_0-rmse:4970.66799      validation_1-rmse:5662.25928
[5]      validation_0-rmse:4397.13963      validation_1-rmse:5106.75982
[6]      validation_0-rmse:4057.24959      validation_1-rmse:4771.82219
[7]      validation_0-rmse:3822.63958      validation_1-rmse:4610.10440
[8]      validation_0-rmse:3645.34812      validation_1-rmse:4541.40784
[9]      validation_0-rmse:3512.55671      validation_1-rmse:4492.49500
[10]     validation_0-rmse:3375.72474      validation_1-rmse:4488.07927
[11]     validation_0-rmse:3311.14597      validation_1-rmse:4474.10073
[12]     validation_0-rmse:3222.20924      validation_1-rmse:4471.14367
[13]     validation_0-rmse:3123.12721      validation_1-rmse:4490.20394
[14]     validation_0-rmse:3050.29756      validation_1-rmse:4498.37645
[15]     validation_0-rmse:3030.27200      validation_1-rmse:4494.12635
[16]     validation_0-rmse:2954.61214      validation_1-rmse:4496.12827
[17]     validation_0-rmse:2892.54924      validation_1-rmse:4521.35732
[18]     validation_0-rmse:2764.14245      validation_1-rmse:4546.57569
[19]     validation_0-rmse:2696.99813      validation_1-rmse:4542.33188
[20]     validation_0-rmse:2632.43704      validation_1-rmse:4530.55994
```

[21]	validation_0-rmse:2584.61584	validation_1-rmse:4535.77625
[22]	validation_0-rmse:2527.93047	validation_1-rmse:4543.95639
[23]	validation_0-rmse:2468.29461	validation_1-rmse:4568.30553
[24]	validation_0-rmse:2403.46871	validation_1-rmse:4583.96935
[25]	validation_0-rmse:2365.15049	validation_1-rmse:4597.30591
[26]	validation_0-rmse:2328.78542	validation_1-rmse:4612.73905
[27]	validation_0-rmse:2276.17594	validation_1-rmse:4621.37732
[28]	validation_0-rmse:2257.53377	validation_1-rmse:4630.06126
[29]	validation_0-rmse:2227.57430	validation_1-rmse:4632.32797
[30]	validation_0-rmse:2158.27133	validation_1-rmse:4633.59277
[31]	validation_0-rmse:2122.02279	validation_1-rmse:4643.16265
[32]	validation_0-rmse:2108.64767	validation_1-rmse:4642.75280
[33]	validation_0-rmse:2056.91867	validation_1-rmse:4671.44408
[34]	validation_0-rmse:2033.70646	validation_1-rmse:4684.71168
[35]	validation_0-rmse:2029.62547	validation_1-rmse:4684.88997
[36]	validation_0-rmse:1994.54510	validation_1-rmse:4693.26534
[37]	validation_0-rmse:1981.56759	validation_1-rmse:4691.73150
[38]	validation_0-rmse:1926.58067	validation_1-rmse:4695.43395
[39]	validation_0-rmse:1874.85411	validation_1-rmse:4717.92004
[40]	validation_0-rmse:1857.27879	validation_1-rmse:4718.50099
[41]	validation_0-rmse:1830.16744	validation_1-rmse:4709.73574
[42]	validation_0-rmse:1767.69701	validation_1-rmse:4724.83488
[43]	validation_0-rmse:1703.13915	validation_1-rmse:4778.27119
[44]	validation_0-rmse:1656.20379	validation_1-rmse:4788.64363
[45]	validation_0-rmse:1609.83477	validation_1-rmse:4777.29533
[46]	validation_0-rmse:1572.18702	validation_1-rmse:4781.85206
[47]	validation_0-rmse:1567.22527	validation_1-rmse:4783.89930
[48]	validation_0-rmse:1539.31888	validation_1-rmse:4808.19374
[49]	validation_0-rmse:1532.71819	validation_1-rmse:4809.19389
[50]	validation_0-rmse:1509.51840	validation_1-rmse:4819.17982
[51]	validation_0-rmse:1470.37402	validation_1-rmse:4818.69505
[52]	validation_0-rmse:1448.02958	validation_1-rmse:4845.66365
[53]	validation_0-rmse:1442.19708	validation_1-rmse:4843.40563
[54]	validation_0-rmse:1393.94653	validation_1-rmse:4846.84045
[55]	validation_0-rmse:1372.76705	validation_1-rmse:4841.93288
[56]	validation_0-rmse:1338.61655	validation_1-rmse:4873.32958
[57]	validation_0-rmse:1299.63161	validation_1-rmse:4886.52445
[58]	validation_0-rmse:1292.29898	validation_1-rmse:4888.67500
[59]	validation_0-rmse:1282.16949	validation_1-rmse:4890.51711
[60]	validation_0-rmse:1273.10161	validation_1-rmse:4890.99547
[61]	validation_0-rmse:1256.99184	validation_1-rmse:4891.56390

[49]: `XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None, colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1, early_stopping_rounds=50, enable_categorical=False, eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise', importance_type=None, interaction_constraints='',`

```
learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,
max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
missing=nan, monotone_constraints='()', n_estimators=1000,
n_jobs=-1, num_parallel_tree=1, predictor='auto', random_state=42,
reg_alpha=0, reg_lambda=1, ...)
```

6 Checking the Importance of Each Feature


```
[50]: feature_importance_df = pd.DataFrame({
      'Feature':input_df.columns,
      'Importance':model.feature_importances_
})

feature_importance_df.sort_values('Importance', ascending=False)
```

```
[50]:
```

	Feature	Importance
4	smoker	0.903173
2	bmi	0.020471
0	age	0.016071
9	bmi_class_healthy weight	0.012697
10	bmi_class_obese	0.007735
3	children	0.007344
6	region_northwest	0.007163
11	bmi_class_overweight	0.006668
8	region_southwest	0.005663
7	region_southeast	0.005130
1	sex	0.004088
5	region_northeast	0.003797
12	bmi_class_underweight	0.000000

```
[51]: sns.barplot(data=feature_importance_df.sort_values('Importance',
↪ascending=False), y='Feature', x='Importance');
```



The smoker feature appears to be the most important one in determining the charges.

To check the best number of trees.

```
[52]: model.best_ntree_limit
```

```
[52]: 13
```

13 is the best number of trees use in the model.

Checking the R^2 score for train and validation sets.

```
[53]: from sklearn.metrics import r2_score

print('The  $R^2$  Score for the Training Set is: {}'.format(r2_score(y_train,
    ↪model.predict(X_train))))
print('The  $R^2$  Score for the Validation Set is: {}'.format(r2_score(y_val,
    ↪model.predict(X_val))))
```

The R^2 Score for the Training Set is: 0.92416713065968

The R^2 Score for the Validation Set is: 0.8912083506265275

Since it ranges from 0 to 1, where 0 is the minimum and 1 the maximum, the model seems to work properly.

Moreover, another way to check the goodness of a model is by performing $NRMSE = RMSE / (y_{max} - y_{min})$ -> <https://www.statology.org/what-is-a-good-rmse/>

The closer to 0 the better.

```
[54]: from sklearn.metrics import mean_squared_error

rmse = mean_squared_error(y_val, model.predict(X_val), squared=False)

nrmse = rmse / (max(target_df) - min(target_df))

print('The RMSE is: {}'.format(rmse))
print('The NRMSE is: {}'.format(nrmse))
```

The RMSE is: 4471.143622095877

The NRMSE is: 0.07136866421921444

Now it is better to check some examples of the difference between the predicted and the actual values.

```
[55]: y_pred = model.predict(X_val)
y_pred[:10]
```

```
[55]: array([ 9457.824 ,  5656.2124, 13328.179 , 37901.95  ,  5033.9526,
        9560.842 , 37780.    ,  2835.1318,  8372.604 , 10332.222 ],
      dtype=float32)
```

```
[56]: y_val[:10]
```

```
[56]: 900      8688.85885
      1064     5708.86700
      1256    11436.73815
      298     38746.35510
      237     4463.20510
      481     9304.70190
      240     38511.62830
      277     2150.46900
      415     7345.72660
      707     10264.44210
      Name: charges, dtype: float64
```

The predictions, except certain cases, do not seem to be so far from reality.

7 Hyperparameter Tuning

In this section are provided some techniques to improve the score of the model.

The hyperparameters can be found here: https://xgboost.readthedocs.io/en/latest/python/python_api.html#xgb


7.1 Max Depth

```
[57]: def optimal_max_depth(number):
      max_depth_model = XGBRegressor(n_jobs=-1, n_estimators=13, max_depth=number,
      ↪random_state=42)
      max_depth_model.fit(X_train, y_train)
      train_rmse = mean_squared_error(y_train, max_depth_model.predict(X_train),
      ↪squared=False)
      val_rmse = mean_squared_error(y_val, max_depth_model.predict(X_val),
      ↪squared=False)
      return {'Max_Depth':number, 'Train_Rmse':train_rmse, 'Validation_Rmse':
      ↪val_rmse}
```

```
[58]: max_depth_df = pd.DataFrame([optimal_max_depth(number) for number in
      ↪range(1,10)]).sort_values('Validation_Rmse')
      max_depth_df.head()
```

```
[58]:   Max_Depth  Train_Rmse  Validation_Rmse
      2         3  4264.770476         4251.996995
      3         4  4038.957467         4280.987639
      4         5  3674.277444         4356.719852
      1         2  4458.462012         4405.776079
      5         6  3222.209225         4471.143622
```

```
[59]: sns.lineplot(data=max_depth_df, x='Max_Depth', y='Train_Rmse')
      sns.lineplot(data=max_depth_df, x='Max_Depth', y='Validation_Rmse')
      plt.ylabel('Train and Validation RMSE');
```



At “max_depth”=3 the validation rmse starts increasing.


7.2 Learning Rate

```
[60]: def optimal_learning_rate(number):
        model = XGBRegressor(n_jobs=-1, n_estimators=13, max_depth=3,
        ↪learning_rate=number, random_state=42)
        model.fit(X_train, y_train)
        train_rmse = mean_squared_error(y_train, model.predict(X_train),
        ↪squared=False)
        val_rmse = mean_squared_error(y_val, model.predict(X_val), squared=False)
        return {'Learning_Rate':number, 'Train_Rmse':train_rmse, 'Validation_Rmse':
        ↪val_rmse}
```

```
[61]: learning_rate_df = pd.DataFrame([optimal_learning_rate(number) for number in np.
        ↪arange(0.1,1,0.05)]).sort_values('Validation_Rmse')
learning_rate_df.head()
```

```
[61]:   Learning_Rate   Train_Rmse  Validation_Rmse
5           0.35  4217.498960        4238.039332
8           0.50  4097.225842        4244.259593
9           0.55  4008.885655        4247.309357
4           0.30  4264.770476        4251.996995
7           0.45  4131.279452        4257.320246
```

```
[62]: sns.lineplot(data=learning_rate_df, x='Learning_Rate', y='Train_Rmse')
sns.lineplot(data=learning_rate_df, x='Learning_Rate', y='Validation_Rmse')
plt.ylabel('Train and Validation RMSE');
```



0.35 is the best value.

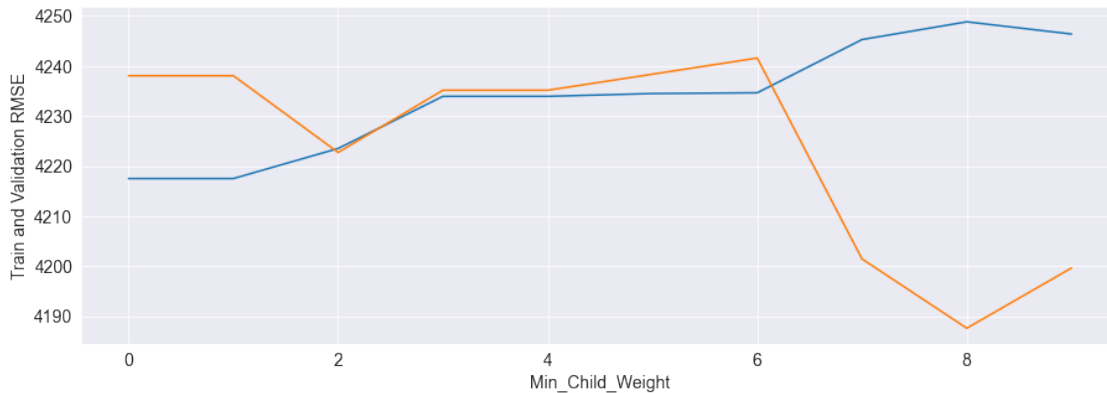
7.3 Min Child Weight

```
[63]: def optimal_child_weight(number):
        model = XGBRegressor(n_jobs=-1, n_estimators=13, max_depth=3,
        ↪ learning_rate=0.35, min_child_weight=number, random_state=42)
        model.fit(X_train, y_train)
        train_rmse = mean_squared_error(y_train, model.predict(X_train),
        ↪ squared=False)
        val_rmse = mean_squared_error(y_val, model.predict(X_val), squared=False)
        return {'Min_Child_Weight': number, 'Train_Rmse': train_rmse,
        ↪ 'Validation_Rmse': val_rmse}
```

```
[64]: child_weight_df = pd.DataFrame([optimal_child_weight(number) for number in
        ↪ range(0,10)]).sort_values('Validation_Rmse')
child_weight_df.head()
```

```
[64]:   Min_Child_Weight  Train_Rmse  Validation_Rmse
8                8  4248.802687      4187.605787
9                9  4246.365003      4199.642400
7                7  4245.259970      4201.440936
2                2  4223.506507      4222.713290
3                3  4233.919525      4235.137942
```

```
[65]: sns.lineplot(data=child_weight_df, x='Min_Child_Weight', y='Train_Rmse')
sns.lineplot(data=child_weight_df, x='Min_Child_Weight', y='Validation_Rmse')
plt.ylabel('Train and Validation RMSE');
```

At level 8 the rmse decreases a lot.

8 Creating a New Model

```
[66]: final_model = XGBRegressor(n_jobs=-1, n_estimators=13, max_depth=3,
    ↪learning_rate=0.35, min_child_weight=8, random_state=42)
final_model.fit(X_train, y_train)
```

```
[66]: XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
    colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
    early_stopping_rounds=None, enable_categorical=False,
    eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',
    importance_type=None, interaction_constraints='',
    learning_rate=0.35, max_bin=256, max_cat_to_onehot=4,
    max_delta_step=0, max_depth=3, max_leaves=0, min_child_weight=8,
    missing=nan, monotone_constraints='()', n_estimators=13, n_jobs=-1,
    num_parallel_tree=1, predictor='auto', random_state=42,
    reg_alpha=0, reg_lambda=1, ...)
```

Checking new scores.

```
[67]: print('The R^2 Score for the Training Set is: {}'.format(r2_score(y_train,
    ↪final_model.predict(X_train))))
print('The R^2 Score for the Validation Set is: {}'.format(r2_score(y_val,
    ↪final_model.predict(X_val))))
```

The R² Score for the Training Set is: 0.8681490868783452

The R² Score for the Validation Set is: 0.9045689059547112

The R² score of the validation set has improved.

```
[68]: rmse = mean_squared_error(y_val, final_model.predict(X_val), squared=False)

nrmse = rmse/(max(target_df)-min(target_df))
```

```
print('The RMSE is: {}'.format(rmse))
print('The NRMSE is: {}'.format(nrmse))
```

The RMSE is: 4187.605787188936
The NRMSE is: 0.0668428161939097

The rmse and the nrmse have decreased, the hyperparameter tuning was completed successfully.

9 Making Predictions on New Inputs

```
[80]: def make_new_predictions(single_input):
    single_input_df = pd.DataFrame([single_input])
    single_input_df['bmi_class']=single_input_df.bmi.map(bmi_estimator)
    single_input_df['sex'] = single_input_df.sex.map(sex_dict)
    single_input_df['smoker'] = single_input_df.smoker.map(smoker_dict)
    single_input_df[encoded_cols] = encoder.
    ↪transform(single_input_df[columns_to_encode])
    single_input_df[columns_to_scale] = scaler.
    ↪transform(single_input_df[columns_to_scale])
    single_input_df = single_input_df.drop(columns=['region', 'bmi_class'])
    pred = final_model.predict(single_input_df)[0]
    return 'The charge is: ${}'.format(pred)
```

```
[70]: new_input = {
    'age':23,
    'sex':'male',
    'bmi':28.1,
    'children':0,
    'smoker':'no',
    'region':'northwest'
}
```

```
[81]: make_new_predictions(new_input)
```

```
[81]: 'The charge is: $4507.5537109375'
```

10 Extra: Using an Artificial Neural Network

```
[ ]: import tensorflow as tf
from tensorflow import keras
```

```
[247]: ann = keras.Sequential([
    keras.layers.Dense(10, input_shape=(13,), activation='relu'),
    keras.layers.Dense(3, activation='relu'),
```

```

        keras.layers.Dense(1, activation='linear')
    ])

    ann.compile(
        optimizer='sgd',
        loss='mse',
        metrics=[tfl.keras.metrics.RootMeanSquaredError()]
    )

    ann.fit(X_train, y_train, epochs=100)

```

```

Epoch 1/100
34/34 [=====] - 0s 2ms/step - loss: 1023015744.0000 -
root_mean_squared_error: 31984.6172
Epoch 2/100
34/34 [=====] - 0s 2ms/step - loss: 176589888.0000 -
root_mean_squared_error: 13288.7129
Epoch 3/100
34/34 [=====] - 0s 2ms/step - loss: 147162272.0000 -
root_mean_squared_error: 12131.0459
Epoch 4/100
34/34 [=====] - 0s 2ms/step - loss: 139552272.0000 -
root_mean_squared_error: 11813.2246
Epoch 5/100
34/34 [=====] - 0s 2ms/step - loss: 137603088.0000 -
root_mean_squared_error: 11730.4346
Epoch 6/100
34/34 [=====] - 0s 2ms/step - loss: 137115984.0000 -
root_mean_squared_error: 11709.6533
Epoch 7/100
34/34 [=====] - 0s 2ms/step - loss: 136999296.0000 -
root_mean_squared_error: 11704.6699
Epoch 8/100
34/34 [=====] - 0s 2ms/step - loss: 136994272.0000 -
root_mean_squared_error: 11704.4551
Epoch 9/100
34/34 [=====] - 0s 2ms/step - loss: 136996112.0000 -
root_mean_squared_error: 11704.5342
Epoch 10/100
34/34 [=====] - 0s 2ms/step - loss: 136972224.0000 -
root_mean_squared_error: 11703.5137
Epoch 11/100
34/34 [=====] - 0s 2ms/step - loss: 137018096.0000 -
root_mean_squared_error: 11705.4727
Epoch 12/100
34/34 [=====] - 0s 2ms/step - loss: 136983792.0000 -
root_mean_squared_error: 11704.0078
Epoch 13/100

```

34/34 [=====] - 0s 2ms/step - loss: 136982960.0000 -
root_mean_squared_error: 11703.9717
Epoch 14/100
34/34 [=====] - 0s 2ms/step - loss: 137011376.0000 -
root_mean_squared_error: 11705.1855
Epoch 15/100
34/34 [=====] - 0s 2ms/step - loss: 137011952.0000 -
root_mean_squared_error: 11705.2109
Epoch 16/100
34/34 [=====] - 0s 2ms/step - loss: 136983872.0000 -
root_mean_squared_error: 11704.0107
Epoch 17/100
34/34 [=====] - 0s 2ms/step - loss: 136980240.0000 -
root_mean_squared_error: 11703.8555
Epoch 18/100
34/34 [=====] - 0s 2ms/step - loss: 136984624.0000 -
root_mean_squared_error: 11704.0430
Epoch 19/100
34/34 [=====] - 0s 2ms/step - loss: 137001376.0000 -
root_mean_squared_error: 11704.7588
Epoch 20/100
34/34 [=====] - 0s 2ms/step - loss: 137019328.0000 -
root_mean_squared_error: 11705.5254
Epoch 21/100
34/34 [=====] - 0s 2ms/step - loss: 136977648.0000 -
root_mean_squared_error: 11703.7451
Epoch 22/100
34/34 [=====] - 0s 2ms/step - loss: 137017584.0000 -
root_mean_squared_error: 11705.4512
Epoch 23/100
34/34 [=====] - 0s 2ms/step - loss: 137005824.0000 -
root_mean_squared_error: 11704.9482
Epoch 24/100
34/34 [=====] - 0s 2ms/step - loss: 137007344.0000 -
root_mean_squared_error: 11705.0137
Epoch 25/100
34/34 [=====] - 0s 2ms/step - loss: 136976592.0000 -
root_mean_squared_error: 11703.7002
Epoch 26/100
34/34 [=====] - 0s 2ms/step - loss: 136965056.0000 -
root_mean_squared_error: 11703.2070
Epoch 27/100
34/34 [=====] - 0s 2ms/step - loss: 137021824.0000 -
root_mean_squared_error: 11705.6318
Epoch 28/100
34/34 [=====] - 0s 2ms/step - loss: 136975072.0000 -
root_mean_squared_error: 11703.6348
Epoch 29/100

34/34 [=====] - 0s 2ms/step - loss: 137020176.0000 -
root_mean_squared_error: 11705.5615
Epoch 30/100
34/34 [=====] - 0s 2ms/step - loss: 137025360.0000 -
root_mean_squared_error: 11705.7832
Epoch 31/100
34/34 [=====] - 0s 2ms/step - loss: 137002784.0000 -
root_mean_squared_error: 11704.8184
Epoch 32/100
34/34 [=====] - 0s 2ms/step - loss: 136961808.0000 -
root_mean_squared_error: 11703.0684
Epoch 33/100
34/34 [=====] - 0s 2ms/step - loss: 136978816.0000 -
root_mean_squared_error: 11703.7949
Epoch 34/100
34/34 [=====] - 0s 1ms/step - loss: 136985920.0000 -
root_mean_squared_error: 11704.0986
Epoch 35/100
34/34 [=====] - 0s 2ms/step - loss: 136969104.0000 -
root_mean_squared_error: 11703.3799
Epoch 36/100
34/34 [=====] - 0s 2ms/step - loss: 136965648.0000 -
root_mean_squared_error: 11703.2324
Epoch 37/100
34/34 [=====] - 0s 2ms/step - loss: 136972256.0000 -
root_mean_squared_error: 11703.5146
Epoch 38/100
34/34 [=====] - 0s 2ms/step - loss: 136984384.0000 -
root_mean_squared_error: 11704.0332
Epoch 39/100
34/34 [=====] - 0s 2ms/step - loss: 136975968.0000 -
root_mean_squared_error: 11703.6729
Epoch 40/100
34/34 [=====] - 0s 2ms/step - loss: 136966528.0000 -
root_mean_squared_error: 11703.2695
Epoch 41/100
34/34 [=====] - 0s 2ms/step - loss: 136998064.0000 -
root_mean_squared_error: 11704.6172
Epoch 42/100
34/34 [=====] - 0s 2ms/step - loss: 137011952.0000 -
root_mean_squared_error: 11705.2109
Epoch 43/100
34/34 [=====] - 0s 2ms/step - loss: 136979376.0000 -
root_mean_squared_error: 11703.8193
Epoch 44/100
34/34 [=====] - 0s 2ms/step - loss: 136968944.0000 -
root_mean_squared_error: 11703.3730
Epoch 45/100

34/34 [=====] - 0s 2ms/step - loss: 137008480.0000 -
root_mean_squared_error: 11705.0625
Epoch 46/100
34/34 [=====] - 0s 2ms/step - loss: 136984288.0000 -
root_mean_squared_error: 11704.0283
Epoch 47/100
34/34 [=====] - 0s 2ms/step - loss: 137016272.0000 -
root_mean_squared_error: 11705.3945
Epoch 48/100
34/34 [=====] - 0s 1ms/step - loss: 136949136.0000 -
root_mean_squared_error: 11702.5273
Epoch 49/100
34/34 [=====] - 0s 1ms/step - loss: 136987040.0000 -
root_mean_squared_error: 11704.1465
Epoch 50/100
34/34 [=====] - 0s 1ms/step - loss: 137002672.0000 -
root_mean_squared_error: 11704.8145
Epoch 51/100
34/34 [=====] - 0s 2ms/step - loss: 137007120.0000 -
root_mean_squared_error: 11705.0039
Epoch 52/100
34/34 [=====] - 0s 1ms/step - loss: 137017168.0000 -
root_mean_squared_error: 11705.4336
Epoch 53/100
34/34 [=====] - 0s 2ms/step - loss: 136956448.0000 -
root_mean_squared_error: 11702.8389
Epoch 54/100
34/34 [=====] - 0s 2ms/step - loss: 137008176.0000 -
root_mean_squared_error: 11705.0488
Epoch 55/100
34/34 [=====] - 0s 1ms/step - loss: 136995856.0000 -
root_mean_squared_error: 11704.5225
Epoch 56/100
34/34 [=====] - 0s 2ms/step - loss: 137006960.0000 -
root_mean_squared_error: 11704.9971
Epoch 57/100
34/34 [=====] - 0s 1ms/step - loss: 136971888.0000 -
root_mean_squared_error: 11703.4990
Epoch 58/100
34/34 [=====] - 0s 2ms/step - loss: 137011424.0000 -
root_mean_squared_error: 11705.1875
Epoch 59/100
34/34 [=====] - 0s 2ms/step - loss: 136971648.0000 -
root_mean_squared_error: 11703.4883
Epoch 60/100
34/34 [=====] - 0s 1ms/step - loss: 136953808.0000 -
root_mean_squared_error: 11702.7266
Epoch 61/100

34/34 [=====] - 0s 1ms/step - loss: 137024688.0000 -
root_mean_squared_error: 11705.7549
Epoch 62/100
34/34 [=====] - 0s 2ms/step - loss: 137001072.0000 -
root_mean_squared_error: 11704.7461
Epoch 63/100
34/34 [=====] - 0s 2ms/step - loss: 137007056.0000 -
root_mean_squared_error: 11705.0010
Epoch 64/100
34/34 [=====] - 0s 2ms/step - loss: 136997936.0000 -
root_mean_squared_error: 11704.6113
Epoch 65/100
34/34 [=====] - 0s 1ms/step - loss: 136983520.0000 -
root_mean_squared_error: 11703.9961
Epoch 66/100
34/34 [=====] - 0s 1ms/step - loss: 136999584.0000 -
root_mean_squared_error: 11704.6826
Epoch 67/100
34/34 [=====] - 0s 2ms/step - loss: 136985328.0000 -
root_mean_squared_error: 11704.0732
Epoch 68/100
34/34 [=====] - 0s 2ms/step - loss: 136983984.0000 -
root_mean_squared_error: 11704.0156
Epoch 69/100
34/34 [=====] - 0s 1ms/step - loss: 137004432.0000 -
root_mean_squared_error: 11704.8896
Epoch 70/100
34/34 [=====] - 0s 1ms/step - loss: 136978064.0000 -
root_mean_squared_error: 11703.7627
Epoch 71/100
34/34 [=====] - 0s 1ms/step - loss: 136978176.0000 -
root_mean_squared_error: 11703.7676
Epoch 72/100
34/34 [=====] - 0s 1ms/step - loss: 136978096.0000 -
root_mean_squared_error: 11703.7646
Epoch 73/100
34/34 [=====] - 0s 1ms/step - loss: 136965952.0000 -
root_mean_squared_error: 11703.2451
Epoch 74/100
34/34 [=====] - 0s 1ms/step - loss: 137002048.0000 -
root_mean_squared_error: 11704.7871
Epoch 75/100
34/34 [=====] - 0s 2ms/step - loss: 136999568.0000 -
root_mean_squared_error: 11704.6816
Epoch 76/100
34/34 [=====] - 0s 1ms/step - loss: 137017312.0000 -
root_mean_squared_error: 11705.4395
Epoch 77/100

34/34 [=====] - 0s 1ms/step - loss: 136952224.0000 -
root_mean_squared_error: 11702.6592
Epoch 78/100
34/34 [=====] - 0s 2ms/step - loss: 136991488.0000 -
root_mean_squared_error: 11704.3359
Epoch 79/100
34/34 [=====] - 0s 2ms/step - loss: 137023936.0000 -
root_mean_squared_error: 11705.7227
Epoch 80/100
34/34 [=====] - 0s 2ms/step - loss: 136962896.0000 -
root_mean_squared_error: 11703.1152
Epoch 81/100
34/34 [=====] - 0s 1ms/step - loss: 136982448.0000 -
root_mean_squared_error: 11703.9502
Epoch 82/100
34/34 [=====] - 0s 2ms/step - loss: 136994624.0000 -
root_mean_squared_error: 11704.4707
Epoch 83/100
34/34 [=====] - 0s 1ms/step - loss: 137003904.0000 -
root_mean_squared_error: 11704.8662
Epoch 84/100
34/34 [=====] - 0s 1ms/step - loss: 136984304.0000 -
root_mean_squared_error: 11704.0293
Epoch 85/100
34/34 [=====] - 0s 2ms/step - loss: 136977696.0000 -
root_mean_squared_error: 11703.7471
Epoch 86/100
34/34 [=====] - 0s 2ms/step - loss: 136982464.0000 -
root_mean_squared_error: 11703.9512
Epoch 87/100
34/34 [=====] - 0s 2ms/step - loss: 137010544.0000 -
root_mean_squared_error: 11705.1504
Epoch 88/100
34/34 [=====] - 0s 2ms/step - loss: 136988256.0000 -
root_mean_squared_error: 11704.1982
Epoch 89/100
34/34 [=====] - 0s 2ms/step - loss: 137008288.0000 -
root_mean_squared_error: 11705.0537
Epoch 90/100
34/34 [=====] - 0s 1ms/step - loss: 136993968.0000 -
root_mean_squared_error: 11704.4424
Epoch 91/100
34/34 [=====] - 0s 2ms/step - loss: 136968336.0000 -
root_mean_squared_error: 11703.3477
Epoch 92/100
34/34 [=====] - 0s 2ms/step - loss: 136985888.0000 -
root_mean_squared_error: 11704.0967
Epoch 93/100


```

34/34 [=====] - 0s 1ms/step - loss: 137019552.0000 -
root_mean_squared_error: 11705.5352
Epoch 94/100
34/34 [=====] - 0s 2ms/step - loss: 137007712.0000 -
root_mean_squared_error: 11705.0293
Epoch 95/100
34/34 [=====] - 0s 2ms/step - loss: 137005808.0000 -
root_mean_squared_error: 11704.9482
Epoch 96/100
34/34 [=====] - 0s 2ms/step - loss: 137000880.0000 -
root_mean_squared_error: 11704.7373
Epoch 97/100
34/34 [=====] - 0s 2ms/step - loss: 137021120.0000 -
root_mean_squared_error: 11705.6025
Epoch 98/100
34/34 [=====] - 0s 2ms/step - loss: 136983904.0000 -
root_mean_squared_error: 11704.0127
Epoch 99/100
34/34 [=====] - 0s 2ms/step - loss: 136972720.0000 -
root_mean_squared_error: 11703.5342
Epoch 100/100
34/34 [=====] - 0s 2ms/step - loss: 136958528.0000 -
root_mean_squared_error: 11702.9277

```

[247]: <keras.callbacks.History at 0x181203b57b0>

It does not seem to perform so well in this case, for sure it needs further implementation.